

DATA266 Lab 1 — LLM Pretraining, Sentiment Classification and CycleGAN Style Transfer

Team PAIR_01 · Khushi Donda · Dipin

Fall 2026

Contents

1. Summary	2
2. Task 1 — GPT-style language model from scratch	2
2.1 Objective, data and preprocessing	2
2.2 Architecture and hyperparameters	2
2.3 Results	3
2.4 Generation quality and failure analysis	3
2.5 Strengths, weaknesses, limitations and next steps	5
3. Task 2 — Yelp Polarity sentiment classification	5
3.1 Data and preprocessing	5
3.2 Models and hyperparameters	5
3.3 Results	6
3.4 Error review (experimental 2)	9
3.5 Strengths, weaknesses, limitations and next steps	9
4. Task 3 — CycleGAN image style transfer (Monet and photos)	10
4.1 Models and hyperparameters	10
4.2 Evaluation protocols	11
4.3 Comparison of results	11
4.4 Human audit (Khushi's model)	12
4.5 Kaggle class competition	13
4.6 Training behaviour	13
4.7 Qualitative results	15
4.8 Failure analysis	15
4.9 Joint analysis	16
5. Cross-task synthesis	16
6. Reproducibility and evidence	18
References	19

GitHub repository: <https://github.com/khushidonda/DATA266-Lab1-Khushi-Dipin>

Team ownership. Khushi led Task 1, Dipin led Task 2, and both team members independently developed and evaluated models for Task 3.

1. Summary

Task	Model(s)	Headline result
1. GPT from scratch (TinyStories, character level)	4-block decoder-only Transformer, 574,830 parameters	validation cross-entropy 0.8643 nats/char, perplexity 2.373, 1.247 bits per character, top-1 accuracy 72.75%
2. Yelp Polarity sentiment	1-layer LSTM, 2-layer BiLSTM, BiLSTM + attention (about 8.2 to 8.8 M parameters each)	test accuracy 0.9518 / 0.9531 / 0.9533, macro-F1 0.9518 / 0.9531 / 0.9533
3. CycleGAN, Monet and photo	Khushi: ResNet-9 CycleGAN, 28.29 M parameters. Dipin: ResNet-9 CycleGAN, 28.27 M parameters, two versions	course evaluator average FID 98.61 (Khushi), 102.39 (Dipin v2), 100.41 (Dipin v1); team leaderboard rank 29, score -49.5121

Every number in this report is read from an artifact in the repository; section 6 maps results to files, checkpoints and SHA-256 values. Statistical statements are stated with their sample sizes, and evaluation protocols are named wherever they differ between models.

2. Task 1 — GPT-style language model from scratch

2.1 Objective, data and preprocessing

The task is next-character prediction with a decoder-only Transformer implemented without any prebuilt Transformer or attention module. The data are 100,000 training and 10,000 validation stories sampled with seed 42 from the 2,119,719-story training split of TinyStories (Eldan and Li, 2023; Hugging Face revision f54c09fd). The tokenizer is a from-scratch character vocabulary of 110 symbols built from the training stories only, with `char_to_idx` and `idx_to_char` dictionaries. Each story is cut into non-overlapping windows of 129 characters (input = first 128, target = the same window shifted by one), giving 644,854 training and 64,576 validation sequences (82.5 M and 8.3 M target characters).

2.2 Architecture and hyperparameters

The model follows the pre-layer-norm GPT design (Vaswani et al., 2017): a trainable token embedding (110 x 128) plus a learned positional embedding (128 x 128), four Transformer blocks and a final layer norm with a linear language-modelling head. Each block applies $x = x + \text{Attn}(\text{LN}(x))$ and $x = x + \text{FFN}(\text{LN}(x))$. Multi-head causal self-attention is written by hand: separate query, key, value and output projections, scaled dot-product scores, a lower-triangular mask applied with `masked_fill(-inf)` before the softmax, and a manual head split and merge. The feed-forward layer is `Linear(128, 256)`, `GELU`, `Linear(256, 128)`.

Item	Value
Sequence length / embedding dim / heads / blocks / feed-forward dim	128 / 128 / 4 (32 per head) / 4 / 256
Dropout	0.1 (embedding, attention, residual, feed-forward)
Parameters	574,830
Optimizer	AdamW, learning rate 3×10^{-4} , weight decay 0.01
Schedule	1,000 linear warm-up steps, then cosine decay to 0 at step 100,760

Item	Value
Batch size / epochs / steps	64 / 10 / 100,760 (10,076 per epoch)
Loss, precision, clipping	cross-entropy, FP32, no gradient clipping
Generation	temperature 0.8, 200 new characters, 5 prompts x 10 samples, seed 42
Hardware / software	NVIDIA GeForce RTX 5090, PyTorch 2.14.0 with CUDA 13.0 (Docker image khushi-task1-gpu:torch2.14-cu130)
Seed	42

2.3 Results

The study evaluates one locked configuration; the final model is the epoch-9 checkpoint, which also has the lowest validation loss (checkpoint epoch_9.pt, SHA-256 5bd9af9b...d398b).

Metric	Validation	Training (eval mode)
Cross-entropy (nats/char)	0.8643	0.8639
Perplexity	2.373	2.372
Bits per character	1.2469	1.2463
Generalization gap (val – train CE)	+0.00040	
Top-1 next-character accuracy	72.75%	72.76%
Distinct-1 / 2 / 3 (character level)	0.0051 / 0.0436 / 0.1592	
Distinct-1 / 2 / 3 (word level, secondary)	0.2395 / 0.6463 / 0.8569	
Repeated 4-gram rate (character / word)	0.1309 / 0.0020	
Gradient norm (all 100,760 steps)	mean 0.408, median 0.413, p99 0.464, max 1.437 (step 126, warm-up)	
NaN / Inf events	0	
Parameters	574,830	
Training throughput / generation throughput	307,817 characters/s / 566.4 characters/s (batch 1)	
Peak GPU memory	444.1 MiB allocated (480.0 MiB reserved)	
Total training time	2,705.1 s (45 min 05 s)	

Epoch	0	1	2	3	4	5	6	7	8	9
Train CE (dropout on)	1.512	1.078	1.017	0.989	0.971	0.959	0.951	0.945	0.942	0.940
Validation CE	1.043	0.954	0.918	0.899	0.885	0.877	0.871	0.866	0.865	0.864

2.4 Generation quality and failure analysis

Samples reproduce the surface form of short children’s stories: common openings, names, punctuation, dialogue and simple declarative sentences. Three failure cases from the 50 generated samples (eval_epoch_9/generations.jsonl) illustrate the main weaknesses.




Figure 1: Training and validation loss per epoch. Training loss is measured with dropout on, so it sits above the eval-mode training loss.




Figure 2: Training dynamics from the step-level log: loss, gradient norm, learning rate and throughput.

Failure type	Generated snippet (prompt in the first words)	Observation
Phrase and template repetition (sample 8)	“Once upon a time, there was a little girl named Lily. She loved to play in the park. One day, a little girl named Lily went to the park. She made a loud noise...”	The phrase “a little girl named Lily” restarts the same subject description instead of advancing the story: a locally probable template is reused.
Semantic incoherence (sample 0)	“...She had a big box to the mountain. One day, her mom saw a big box to sparkle and the spider. They also learned that her name cleaning the boy...”	Syntax and rhythm of a children’s story are kept, but phrases such as “a big box to sparkle” do not form coherent relations.
Narrative and dialogue coherence (sample 46)	“Tom was sad because he was too high and felt the wallets and the friends laughed inside... said, ”We will write as me fluffer?””	Dialogue formatting is preserved while meaning and entity continuity deteriorate.

2.5 Strengths, weaknesses, limitations and next steps

Strengths. Optimization is stable over all 100,760 steps (no NaN, maximum gradient norm 1.44 at warm-up) and the evaluation-mode train and validation losses differ by only 0.0004 nats, so there is little evidence of overfitting on this split. A compact 0.57 M-parameter model learns strong local character structure.

Weaknesses and limitations. Local fluency is much stronger than long-range coherence: with a 128-character context the model cannot track entities across sentences, and character-level sampling can produce plausible but invalid words (the word-level unseen share is 1.1%). The character-level repeated 4-gram rate (0.131) reflects both genuine template reuse and the small alphabet. One configuration was evaluated, so the effect of context length and capacity is not measured here.

Next steps. Hold the evaluation protocol fixed and test longer context windows and larger capacity, and decoding controls that penalize repeated n-grams; treat each as a separate experiment so the reference run stays unchanged.

3. Task 2 — Yelp Polarity sentiment classification

3.1 Data and preprocessing

The data are fancyzhx/yelp_polarity: 560,000 training reviews and 38,000 official test reviews, with no null entries, no duplicates and no invalid labels. A stratified 90/10 split of the training pool (seed 42) gives 504,000 training and 56,000 validation reviews; the test set is untouched. All splits are exactly balanced (50/50). Raw reviews are longer for the negative class (mean 151.5 words, median 112) than for the positive class (mean 114.5, median 84).

Preprocessing runs in this order: unescape literal `\n`, lowercase, expand negation contractions, strip punctuation, whitespace tokenization, NLTK English stopword removal with 11 negation words kept (no, not, nor, never, none, nobody, nothing, neither, nowhere, cannot, without), and WordNet lemmatization. The vocabulary is built on the training split only (minimum frequency 3, cap 40,000, plus `<PAD>` and `<UNK>`, 40,002 entries); sequences are cut to 200 tokens by keeping the first 150 and last 50 (4.46% of test reviews are truncated; the test unknown-token rate is 0.70%). Embeddings are learned from scratch; no pretrained vectors or language models are used.

3.2 Models and hyperparameters

Three models share the pipeline and differ only in the architecture:




Figure 3: Class balance and review length distribution.

Model	Architecture	Parameters
Baseline	Embedding(V, 200), 1-layer unidirectional LSTM(128), hidden state at the true sequence length, dropout 0.3, Linear(128, 1)	8,169,489
Experimental 1	Embedding(V, 200), 2-layer BiLSTM (128 per direction, inter-layer dropout 0.3), concatenated final forward and backward states, dropout, Linear(256, 1)	8,733,841
Experimental 2	The same 2-layer BiLSTM with masked additive (Bahdanau) attention pooling over non-padding steps (attention size 128), dropout, Linear(256, 1)	8,766,865

Shared settings: embedding dimension 200, AdamW (learning rate 2×10^{-3} , weight decay 10^{-4}), 5% warm-up then cosine decay, batch size 256, at most 4 epochs with early stopping on validation macro-F1 (patience 2; best epoch 3 for all three), BCE-with-logits loss, embedding dropout 0.2, dropout 0.3, gradient clipping 1.0, seed 42, threshold 0.5. Hardware: NVIDIA GeForce RTX 5090 and Intel Core Ultra 9 285K, PyTorch 2.11.0 with CUDA 12.8.

3.3 Results

Metric (38,000 test reviews)	Baseline	Experimental 1	Experimental 2
Accuracy (= micro precision, recall and F1)	0.9518	0.9531	0.9533
Precision, macro (= weighted)	0.9518	0.9531	0.9533
Recall, macro (= weighted)	0.9518	0.9531	0.9533
F1, macro (= weighted)	0.9518	0.9531	0.9533
ROC-AUC	0.9890	0.9901	0.9903
PR-AUC	0.9889	0.9902	0.9903

Metric (38,000 test reviews)	Baseline	Experimental 1	Experimental 2
Matthews correlation (MCC)	0.9036	0.9062	0.9066
Brier score	0.0376	0.0364	0.0360
Expected calibration error (10 bins)	0.0149	0.0157	0.0144
95% CI, accuracy	[0.9496, 0.9539]	[0.9511, 0.9551]	[0.9512, 0.9554]
95% CI, macro-F1	[0.9496, 0.9539]	[0.9511, 0.9551]	[0.9512, 0.9554]
95% CI, MCC	[0.8993, 0.9078]	[0.9021, 0.9102]	[0.9023, 0.9108]
Confusion matrix (TN / FP / FN / TP)	18,064 / 936 / 896 / 18,104	18,150 / 850 / 932 / 18,068	18,167 / 833 / 941 / 18,059
Parameters	8,169,489	8,733,841	8,766,865
Training time (s) / training examples per second	167.4 / 12,046	543.1 / 3,712	693.3 / 2,908
Inference examples per second	86,703	32,300	27,278
Peak GPU memory (MiB)	577	979	980
Peak host memory (GiB)	3.66	3.68	3.79

Confidence intervals are from 1,000 bootstrap resamples of the test set. Checkpoints: task2_sentiment/dipin/checkpoints/baseline.pt, experimental_1.pt and experimental_2.pt (SHA-256 e6cba920..., d1fbe654..., c3f4da85...).




Figure 4: Training curves of the three models.

Paired McNemar test (exact binomial test on the discordant pairs, 38,000 test reviews):

Comparison	A right, B wrong	A wrong, B right	p-value
Baseline vs Experimental 1	373	423	0.0824
Baseline vs Experimental 2	430	488	0.0599
Experimental 1 vs Experimental 2	431	439	0.8124

None of the three comparisons is statistically significant at the conventional 0.05 level; the baseline versus experimental 2 comparison is borderline. The accuracy differences (0.0013 and 0.0015) are about the size of




Figure 5: ROC, precision-recall and calibration curves.




Figure 6: Confusion matrices: baseline, experimental 1, experimental 2.

the bootstrap interval half-width (about 0.002), so the three models should be regarded as close in accuracy, with the BiLSTM models somewhat better in probability quality (Brier score, AUC).

Robustness by data slice (macro-F1 / error rate):

Slice (n)	Baseline	Experimental 1	Experimental 2
all (38,000)	0.9518 / 4.82%	0.9531 / 4.69%	0.9533 / 4.67%
short reviews, 65 words or fewer (12,776)	0.9488 / 4.97%	0.9506 / 4.80%	0.9492 / 4.93%
medium, 65 to 142 words (12,649)	0.9529 / 4.71%	0.9534 / 4.66%	0.9553 / 4.47%
long, over 142 words (12,575)	0.9507 / 4.78%	0.9523 / 4.61%	0.9524 / 4.60%
contains a negation (28,526)	0.9494 / 4.91%	0.9506 / 4.79%	0.9515 / 4.70%
contains a contrast word (22,551)	0.9467 / 5.28%	0.9484 / 5.11%	0.9483 / 5.12%
truncated, over 200 tokens (1,694)	0.9303 / 6.14%	0.9361 / 5.61%	0.9250 / 6.61%
high unknown-token rate, over 10% (231)	0.8442 / 13.85%	0.8421 / 13.85%	0.8333 / 14.72%

Experimental 2 has the best overall metrics but the highest error on the truncated and high-unknown slices, so its overall advantage does not extend to long or non-English reviews.

3.4 Error review (experimental 2)

Experimental 2 misclassifies 1,774 of 38,000 reviews (4.67%). Twenty errors were selected automatically (not by hand): the 5 most confident false positives, the 5 most confident false negatives, the 5 errors closest to $P = 0.5$ and the 5 most confident errors in the worst slice (high unknown rate). Each was labelled with an error type (task2_sentiment/dipin/failure_analysis.md, candidates in outputs/error_review/error_candidates_experimental_2.csv).

Error type	Count	Example (test index)
Mixed review (praise and complaints in one review)	7	#5185, “But like Clay I did not love the ones that I got ... It is definitely a worthwhile destination for a pancake lover.” True label negative, $P(\text{pos}) = 0.9995$
Unknown words (non-English text, links, broken accents)	5	slice errors #1256, #36104, #35963, #19328, #27150, all with $P(\text{pos}) > 0.99$
Sentiment about something else	3	strong words aimed at other reviewers, earlier owners or other customers
Label noise	2	#29330, “Wow love the place and everything is very clean and new!” labelled negative
Indirect negative	2	#29923, “I would recommend just about every other froyo joint than this one.”
Too little signal	1	very short neutral text near $P = 0.5$

Patterns: most errors are hard rather than random (7 of 20 are mixed reviews, and 4 of the 5 near-threshold errors are mixed); some are label problems (#29330 and one more); and the stopword list removes contrast words that carry meaning (*but*, *than*, *other*, *now*, *before*): reviews that contain “but” have a 5.16% error rate ($n = 20,227$) against 4.11% for the rest ($n = 17,773$). The accented characters stored as \uXXXX codes are damaged by the cleaner for 879 test reviews (2.3%), whose error rate is 7.39%.

Testable fix. Remove *but*, *than*, *other*, *now*, *before*, *only*, *too* from the stopword list (as negations already are), retrain experimental 2 with the same configuration and seed, and compare on the same 38,000 reviews: the error rate on “but” reviews should fall, the error rate on the others should not rise, and a paired McNemar test of old versus new should give $p < 0.05$. A second, smaller fix is to decode the \uXXXX codes before cleaning and replace links by a single <url> token. Two of the 20 errors are label noise that no model change can fix.

3.5 Strengths, weaknesses, limitations and next steps

Strengths. All three models reach about 95% accuracy and AUC near 0.99 with embeddings learned from scratch, and calibration is good (ECE about 0.015). Experimental 2 is best on every aggregate metric.

Weaknesses and limitations. The gains from depth, bidirectionality and attention are small and not significant by McNemar’s test; they cost 3.3 to 4.1 times the training time. Training used at most 4 epochs, with the best validation epoch at 3 for every model. Errors concentrate in mixed, non-English and truncated reviews, and part of the problem is caused by the preprocessing choices above.

Next steps. The stopword fix above, a language filter or character-level fallback for non-English text, and longer-context or hierarchical pooling for the truncated slice.

4. Task 3 — CycleGAN image style transfer (Monet and photos)

Both team members independently implemented a CycleGAN (Zhu et al., 2017) in PyTorch with two ResNet-9 generators and two 70x70 PatchGAN discriminators trained with an LSGAN objective (Mao et al., 2017) on the Kaggle “I’m Something of a Painter Myself” images (300 Monet paintings, 7,038 photographs, 256 x 256 RGB, not paired). No pretrained network generates or alters any translated image; pretrained Inception-v3 and AlexNet are used only to measure images. To avoid confusion between the two code bases, this section names directions by content: **photo to Monet** and **Monet to photo**.

4.1 Models and hyperparameters

	Khushi	Dipin v2 (final)	Dipin v1 (baseline)
Generators	ResNet-9, reflection padding, InstanceNorm, ReLU, Tanh; transposed-convolution upsampling	same, with nearest-neighbour upsampling plus 3x3 convolution	transposed-convolution upsampling
Discriminators	70x70 PatchGAN, InstanceNorm, raw output	same, plus DiffAugment and gradient-norm clipping at 15	70x70 PatchGAN
Parameters	28,285,832	28,273,544	28,273,544
Loss weights	cycle 10, identity 5	cycle 10, identity 1	cycle 10, identity 5
Optimizer	Adam, lr 2×10^{-4} , betas (0.5, 0.999)	same	same
Batch size	1	4	4
Schedule	constant lr for epochs 1 to 100, then linear decay (200 planned); stopped after epoch 40, selected completed epoch 29	50 epochs: 25 constant, 25 linear decay	same
Steps	204,102 at the selected checkpoint (7,038 per epoch)	81,750 (1,635 per epoch)	81,750
Replay pool	50 per domain	50	50
Augmentation	resize 286, random crop 256, random flip	same	same
Precision	FP32	bf16 autocast for convolutions	bf16 autocast
Seed	42	266	266
Hardware	NVIDIA GeForce RTX 4090	NVIDIA GeForce RTX 4090	NVIDIA GeForce RTX 4090
Checkpoint	epoch_28.pt (completed epoch 29), SHA-256 d3124ab2...c89d3c	checkpoints/v2/G_AB.pt, G_BA.pt, D_A.pt, D_B.pt (epoch 50)	checkpoints/*.pt (epoch 50)

Checkpoint file epoch_28.pt corresponds to completed epoch 29 under zero-based checkpoint naming. Khushi’s final checkpoint was selected from the completed epochs 25 to 40 by scoring every checkpoint (JPEG quality 95) with the course evaluator: completed epoch 29 had the lowest average FID (98.61; neighbouring epochs ranged from 99.1 to 103.0) and the lowest average MiFID. A lossless PNG variant of the same checkpoint scored worse (average FID 101.02), so JPEG output was kept. This selection used the same reference images as the final score, so it is model selection on the evaluation set. Dipin’s final model is the improved version v2 (lower

identity weight, DiffAugment on the discriminator inputs, discriminator gradient clipping, resize-convolution upsampling), submitted to the class evaluator; v1 is the baseline it was compared against.

4.2 Evaluation protocols

	Khushi	Dipin
Set	Fixed in-domain evaluation set: all 300 sorted Monet images and the first 300 sorted photos. The generators were trained on these image domains, so the numbers measure in-domain translation quality, not held-out generalization	Held-out inputs: 500 test photos and 30 test Monet paintings never shown to the model (split seed 266; 6,538 / 270 images used for training); distribution metrics compare against all 300 real Monet paintings or the 500 test photos
Feature extractor	Inception-v3 pool features	Inception-v3 pool features

Because the evaluation sets differ, the per-metric values from the two members’ own protocols are shown side by side but are not strictly comparable. The course evaluator is the one measurement taken with identical code and the same N_EVAL, and it is reported in the first block of the table (its inputs are still different images for each member).

4.3 Comparison of results

Metric	Khushi	Dipin v2 (final)	Dipin v1 (baseline)
Course evaluator (identical code, N_EVAL = 300, JPEG)			
FID, photo to Monet	102.219	99.902	102.428
FID, Monet to photo	95.009	104.869	98.383
FID, average of both directions	98.614	102.386	100.405
MiFID, photo to Monet / Monet to photo	0.4069 / 0.4131	0.3913 / 0.4118	0.3976 / 0.4131
MiFID, average	0.4100	0.4016	0.4054
Class metric −(FID + MiFID)/2	−49.512	−51.394	−50.405
Each member’s own evaluation protocol (Khushi: in-domain 300 + 300; Dipin: held-out photos, section 4.2)			
FID, photo to Monet	102.219	92.033	94.907
FID, Monet to photo	95.009	95.493	88.108
KID, photo to Monet	0.00906 ± 0.00179	0.01751 ± 0.00195	0.01430 ± 0.00210
KID, Monet to photo	0.01397 ± 0.00217	0.02598 ± 0.00322	0.01670 ± 0.00236
Precision / recall, photo to Monet	0.400 / 0.613	0.428 / 0.497	0.364 / 0.550
Precision / recall, Monet to photo	0.717 / 0.390	0.597 / 0.292	0.670 / 0.338

Metric	Khushi	Dipin v2 (final)	Dipin v1 (baseline)
Density / coverage, photo to Monet	0.392 / 0.717	0.450 / 0.787	0.415 / 0.807
Density / coverage, Monet to photo	1.289 / 0.940	0.734 / 0.630	0.845 / 0.720
Cycle L1, photo cycle (photo, Monet, photo)	0.0513	0.0497	0.0462
Cycle L1, Monet cycle (Monet, photo, Monet)	0.0421	0.0612	0.0603
LPIPS input vs cycle reconstruction, photo cycle	0.191	0.258	0.158
LPIPS input vs cycle reconstruction, Monet cycle	0.242	0.381	0.232
Content cosine, photo to Monet	0.771	0.752	0.757
Content cosine, Monet to photo	0.796	0.801	0.831
Training			
Parameters	28,285,832	28,273,544	28,273,544
Training time	5.82 h (29 epochs)	3.44 h (50 epochs)	2.61 h (50 epochs)
Throughput (paired photo + Monet samples per second)	9.71	26.39	34.85
Peak GPU memory	2.70 GiB	6.60 GB	4.88 GB
NaN / Inf events	0	0	0
Final-epoch generator loss / mean discriminator loss	3.402 / 0.147	2.645 / 0.142	2.943 / 0.112

Notes. KID is an unbiased MMD² with a polynomial kernel (degree 3), mean $\pm$ standard deviation over 100 subsets of 100 images (Khushi: seed 42, 300 vs 300 images; Dipin: the held-out protocol). Precision and recall use k-nearest-neighbour manifolds with $k = 3$ (Kynkaanniemi et al., 2019); density and coverage use $k = 5$ (Naeem et al., 2020). Cycle L1 is on a $[0, 1]$ pixel scale, computed on raw tensors before JPEG encoding. LPIPS uses the AlexNet backbone (Zhang et al., 2018). Khushi’s final-epoch loss values are for completed epoch 29. The Kaggle-style class metric is calculated from the course evaluator’s FID and MiFID values for each model.

Khushi, additional measurements. LPIPS between input and direct translation (supplementary, since style change is intended): 0.368 (Monet to photo) and 0.390 (photo to Monet). Cycle L1 overall 0.0467 and LPIPS cycle overall 0.2165.

4.4 Human audit (Khushi’s model)

Thirty samples (15 Monet to photo, 15 photo to Monet) were drawn at random with seed 42 from the fixed evaluation set, shuffled, and shown as blinded input | translation panels. Two raters scored each sample independently on a 1 to 5 scale for style quality, content preservation and artifact severity (1 = no visible artifacts, 5 = severe).

Criterion	Rater 1	Rater 2	Combined (mean $\pm$ SD)	Quadratic- weighted kappa	Exact agreement	Mean absolute difference
Style quality (higher is better)	4.00 $\pm$ 0.91	2.97 $\pm$ 0.85	3.48 $\pm$ 1.02	0.130	33.3%	1.10
Content preservation (higher is better)	4.57 $\pm$ 0.57	4.07 $\pm$ 0.94	4.32 $\pm$ 0.81	0.181	50.0%	0.70
Artifact severity (lower is better)	1.60 $\pm$ 0.62	2.43 $\pm$ 1.04	2.02 $\pm$ 0.95	-0.025	36.7%	1.03

Both raters judged content preservation better than style conversion. Inter-rater agreement is low (kappa 0.13, 0.18 and -0.03), so the combined means blend two rating scales and should not be read as a precise consensus. The raters differed most on photo-to-Monet artifacts (rater 1: 1.47, rater 2: 3.20 averaged over the 15 samples) and agreed on Monet to photo (1.73 and 1.67).

4.5 Kaggle class competition

Khushi’s checkpoint `epoch_28.pt` produced the team submission `task3_gan/khushi/outputs/submission.csv` (ID, FID, MiFID / 1, 98.61422779159997, 0.40998050570487976), for which $-(\text{FID} + \text{MiFID})/2 = -49.5121$. The leaderboard screenshot shows team **PairProgramming_Team_1** at **rank 29** with score **-49.5121** (banner: “improving on your previous best of -49.8366”). The rank is the leaderboard position at the time of the screenshot and can change as other teams submit.




Figure 7: Leaderboard row of the team.

4.6 Training behaviour

Khushi. The generator, cycle and identity losses fall smoothly and the gradient norms level off (generator about 20, discriminators about 5 to 6) with no NaN events. The Monet discriminator loss fell from 0.229 to about 0.082 at epoch 26 and then rose to 0.116 by epoch 29, i.e. D_A became somewhat less confident over the last three epochs while D_B stayed near 0.177. This is a moderate change, not evidence of instability by itself, and it is one reason the neighbouring checkpoints were scored.

Dipin. The raw logs show the effect of the v2 changes. In v1 the Monet discriminator is briefly overpowered (its real-minus-fake score gap fell to 0.012), while v2 with gradient clipping and DiffAugment kept a larger minimum gap (0.102) and recorded no NaN steps in either run. v2 has a lower generator loss (2.645 against 2.943) but a higher unweighted cycle loss (0.136 against 0.126) and a larger identity loss, consistent with the lower identity weight.




Figure 8: Khushi: generator loss and discriminator losses per epoch, completed epochs 1 to 29 (the dashed line marks the selected epoch).




Figure 9: Khushi: cycle-consistency and identity losses (unweighted L1).




Figure 10: Khushi: mean gradient norms of the generators and discriminators.




Figure 11: Dipin: epoch-level training curves of v1 (baseline) and v2 (final) from the raw logs.

4.7 Qualitative results

Khushi’s model converts photographs into soft, painterly images with Monet-like palettes while keeping layout, and turns Monet paintings into photograph-like images that keep the scene structure. Dipin’s v2 gives a stronger painted look with pointillist texture and pastel palettes in the photo-to-Monet direction and shows a small yellow mark in the top-left corner of generated Monet images, an artifact documented in Dipin’s analysis.

4.8 Failure analysis

Khushi (task3_gan/khushi/failure_analysis.md; panels from the blinded audit set):

- *Periodic texture over smooth regions (photo to Monet; S09, S14, S27)*. A fine grid or stripe pattern covers skies and clouds; in S14 the night sky is replaced by streaks while the snow foreground is kept. The generator uses transposed-convolution upsampling, which is known to produce such patterns (Odena et al., 2016); this is a hypothesis that was not tested.
- *Loss of detail and contrast (S11, S20)*. The willow in S11 (Monet to photo) becomes darker and blurrier, and the saturated orange terrace in S20 (photo to Monet) washes out to a pale pastel.
- *Partial style conversion (S18, Monet to photo)*. An abstract, heavily textured painting is translated into an output close to the input; content preservation is rated 4 and 5 and style conversion 2 and 1.
- *Rater disagreement on artifacts (above)* means that artifact scores for photo to Monet differ by 1.7 points between raters.

Dipin (task3_gan/dipin/failure_analysis.md, v1 held-out photos selected automatically by lowest content cosine, highest cycle L1 and highest LPIPS). Low-texture scenes such as salt flats and plain skies have the highest cycle error (up to 0.242 against a median of 0.042) and receive hallucinated texture; saturated sunsets change colour yet reconstruct almost perfectly, showing that a low cycle loss alone does not prove a faithful translation (the “steganography” effect of Chu et al., 2017); letterbox borders and flat-sky streak or checkerboard patterns are further failure modes. The v2 run removed the checkerboard grids but introduced blob artifacts and the fixed yellow corner mark.



Figure 12: Khushi, completed epoch 29, first eight sorted inputs. Rows: Monet paintings; Monet to photo; photographs; photo to Monet.

4.9 Joint analysis

Strengths. Both models learn the two mappings and preserve scene content (content cosine 0.75 to 0.83, cycle L1 0.04 to 0.06) without any NaN event. Khushi’s model has the best course-evaluator result (average FID 98.61, class metric -49.512) among the three, and the lowest KID in both directions on its fixed in-domain set; Dipin’s v2 has the best photo-to-Monet FID under the course evaluator (99.90) and on the held-out photos used in Dipin’s protocol (92.03) and the lowest average MiFID (0.4016), at higher training cost than v1.

Weaknesses and limitations. The two sets of per-metric numbers use different evaluation sets (in-domain 300 + 300 versus held-out photos), so the ranking by FID, KID, precision and recall between members must be read with that difference in mind; for Khushi the checkpoint was also selected on the evaluation set. With 300 images per set FID is biased and the KID subset standard deviation (0.002 to 0.003) is not small relative to the differences between directions. In Khushi’s model FID and KID rank the directions differently (FID lower for Monet to photo, KID lower for photo to Monet), and precision and recall are asymmetric (0.717 / 0.390 for Monet to photo, 0.400 / 0.613 for photo to Monet); this is a descriptive reading of 300-image comparisons and does not establish a mode-coverage difference. Khushi’s training stopped at epoch 40 of a planned 200 and never reached the learning-rate decay phase. The human audit has two raters, 30 samples and low agreement.

Next steps. Evaluate all models on a common held-out set with the same code; replace transposed convolutions by resize-convolutions in Khushi’s generator while keeping a smoothness or perceptual term; use DiffAugment on the Monet discriminator only (the photo critic has 6,538 real images) with a moderate identity weight; and extend Khushi’s training into the decay phase.

5. Cross-task synthesis

Across the three tasks the same pattern appears: aggregate metrics improve with model capacity or training tricks, but the gains are small relative to their measurement uncertainty, and the informative evidence comes from failure analysis. In Task 1 evaluation-mode cross-entropy shows no overfitting while generated text still repeats templates and loses coherence. In Task 2 a McNemar test shows the 0.15 point accuracy gain of the BiLSTM with attention is not significant ($p = 0.060$ to 0.082 against the baseline), and the errors are dominated



Figure 13: Dipin v2 (final), fixed grid at epoch 50. Rows: photographs, photo to Monet, cycle reconstruction, Monet paintings, Monet to photo, cycle reconstruction.




Figure 14: Audit panels S09, S14 (top) and S11, S18 (bottom); each panel shows the input on the left and the translation on the right.

by mixed reviews, label noise and preprocessing choices. In Task 3 distribution metrics computed on 300 images have material uncertainty, differ between FID and KID, and cannot replace inspection of artifacts and the human audit. For all tasks the repository stores checkpoints, raw logs, manifests with SHA-256 values and the executed notebooks, so every table above can be regenerated or checked (section 6).

6. Reproducibility and evidence

Result	Evidence
Task 1 metrics, curves, samples	task1_llm/khushi/outputs/task1_full_20260925_204011/ (eval_epoch_9/metrics.json, curves/), task1_llm/khushi/metrics_report.csv, raw log reproducibility/raw_logs/khushi/task1_full_20260925_204011.jsonl, checkpoint task1_llm/khushi/checkpoints/task1_full_20260925_204011/epoch_9.pt, manifest reproducibility/manifests/khushi/manifest.md
Task 2 metrics, plots, errors	task2_sentiment/dipin/metrics_report.csv, outputs/run_summary.json, outputs/plots/, outputs/error_review/, raw logs reproducibility/raw_logs/dipin/task2_*.jsonl, manifest reproducibility/manifests/dipin/task2_manifest.json, checkpoints task2_sentiment/dipin/checkpoints/

Result	Evidence
Task 3, Khushi	task3_gan/khushi/full_metrics_report.csv, results.md, outputs/official_eval/, outputs/final_eval_epoch029/, outputs/human_audit/, outputs/kaggle/, raw logs reproducibility/raw_logs/khushi/task3_khushi_ prod_*, checkpoint-to-result chain reproducibility/ manifests/khushi/task3_checkpoint_result_map.md
Task 3, Dipin	task3_gan/dipin/metrics_report.csv, results.md, outputs/eval/, outputs/v2/, raw logs task3_gan/dipin/outputs/logs/ and outputs/v2/logs/, checkpoints task3_gan/dipin/checkpoints/

The executed notebooks are task1_llm/khushi/task1_khushi.ipynb, task2_sentiment/dipin/src/task2_dipin.ipynb, task3_gan/khushi/src/task3_khushi.ipynb and task3_gan/dipin/src/task3_dipin.ipynb. A single command, `python smoke_test.py` from the repository root, loads every final checkpoint with its own model code and checks a forward pass; the README lists setup, dataset and evaluation commands.

References

1. Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L. and Polosukhin, I. (2017). Attention Is All You Need. *Advances in Neural Information Processing Systems* 30.
2. Eldan, R. and Li, Y. (2023). TinyStories: How Small Can Language Models Be and Still Speak Coherent English? *arXiv:2305.07759*.
3. Zhu, J.-Y., Park, T., Isola, P. and Efros, A. A. (2017). Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks. *IEEE International Conference on Computer Vision (ICCV)*.
4. Mao, X., Li, Q., Xie, H., Lau, R. Y. K., Wang, Z. and Smolley, S. P. (2017). Least Squares Generative Adversarial Networks. *ICCV*.
5. Heusel, M., Ramsauer, H., Unterthiner, T., Nessler, B. and Hochreiter, S. (2017). GANs Trained by a Two Time-Scale Update Rule Converge to a Local Nash Equilibrium. *NeurIPS*.
6. Bińkowski, M., Sutherland, D. J., Arbel, M. and Gretton, A. (2018). Demystifying MMD GANs. *ICLR*.
7. Kynkäänniemi, T., Karras, T., Laine, S., Lehtinen, J. and Aila, T. (2019). Improved Precision and Recall Metric for Assessing Generative Models. *NeurIPS*.
8. Naeem, M. F., Oh, S. J., Uh, Y., Choi, Y. and Yoo, J. (2020). Reliable Fidelity and Diversity Metrics for Generative Models. *ICML*.
9. Zhang, R., Isola, P., Efros, A. A., Shechtman, E. and Wang, O. (2018). The Unreasonable Effectiveness of Deep Features as a Perceptual Metric. *CVPR*.
10. Odena, A., Dumoulin, V. and Olah, C. (2016). Deconvolution and Checkerboard Artifacts. *Distill*.
11. Chu, C., Zhmoginov, A. and Sandler, M. (2017). CycleGAN, a Master of Steganography. *arXiv:1712.02950*.
12. Zhang, X., Zhao, J. and LeCun, Y. (2015). Character-level Convolutional Networks for Text Classification. *NeurIPS (Yelp Polarity)*.
13. Bahdanau, D., Cho, K. and Bengio, Y. (2015). Neural Machine Translation by Jointly Learning to Align and Translate. *ICLR*.
14. Hochreiter, S. and Schmidhuber, J. (1997). Long Short-Term Memory. *Neural Computation* 9(8).
15. McNemar, Q. (1947). Note on the sampling error of the difference between correlated proportions or percentages. *Psychometrika* 12.
16. Cohen, J. (1968). Weighted kappa: nominal scale agreement with provision for scaled disagreement or partial credit. *Psychological Bulletin* 70.